

# Multithreading

## Chapter 30

# Introduction

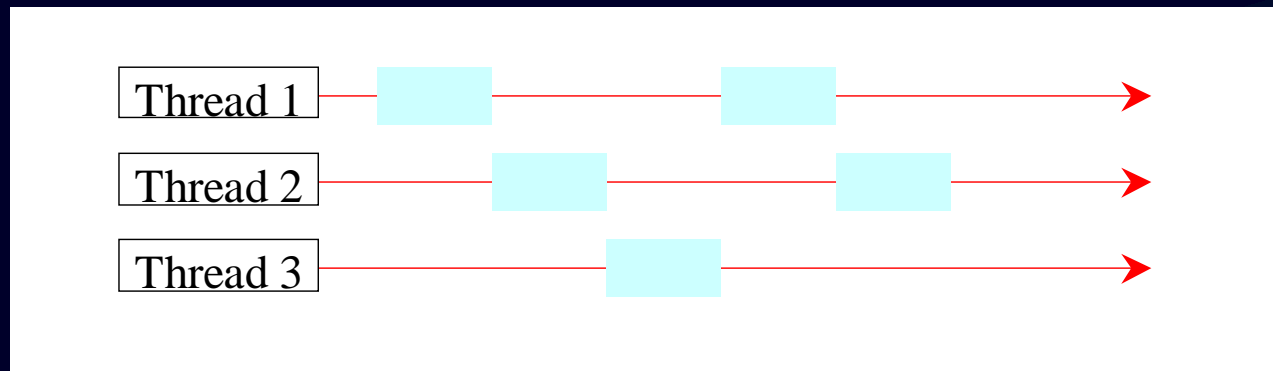
**Multithreading:** The ability for multiple tasks in a program to be executed concurrently.

**Thread:** A program may consist of many tasks that can run concurrently. A **thread** is the flow of execution, from beginning to end, of a task.

# Threads Concept

In single-processor systems, as shown in the figure below, the multiple threads share CPU time, known as *time sharing*, and the operating system is responsible for scheduling and allocating resources to them.


Multiple  
threads  
sharing a  
single CPU



# Threads Concept

On the other hand, threads can be executed simultaneously in multiprocessor systems.

Multiple  
threads on  
multiple  
CPUs



# Threads Concept

- When your program executes as an application, the Java interpreter starts a thread for the **main** method.
- You can create additional threads to run concurrent tasks in the program.
- In Java, each task is an instance of the **Runnable** interface, also called a *runnable object*.
- A *thread* is essentially an object that facilitates the execution of a task.

# Creating Tasks and Threads

- *In order to create a thread, you must create a **task class** and a **thread object**.*
- *A **task class** must implement the **Runnable** interface.*
- *A task must be run from a **thread**.*
- *The **Runnable** interface should be implemented by any class whose instances are intended to be executed by a thread.*
- *The class must define a method of no arguments called **run**.*
- *When an object implementing interface **Runnable** is used to create a thread, starting the thread causes the object's run method to be **called in that separately executing thread**.*
- *The general contract of the method run is that it may take any action whatsoever.*

# Creating Tasks and Threads

`java.lang.Runnable`

TaskClass



```
// Custom task class
public class TaskClass implements Runnable {
    ...
    public TaskClass(...) {
        ...
    }
    // Implement the run method in Runnable
    public void run() {
        // Tell system how to run custom thread
        ...
    }
    ...
}
```

```
// Client class
public class Client {
    ...
    public void someMethod() {
        ...
        // Create an instance of TaskClass
        TaskClass task = new TaskClass(...);

        // Create a thread
        Thread thread = new Thread(task);

        // Start a thread
        thread.start();
        ...
    }
    ...
}
```

# Example: Using the Runnable Interface to Create and Launch Threads (1 of 3)

```
1  public class TaskThreadDemo {
2      public static void main(String[] args) {
3          // Create tasks
4          Runnable printA = new PrintChar('a', 100);
5          Runnable printB = new PrintChar('b', 100);
6          Runnable print100 = new PrintNum(100);
7
8          // Create threads
9          Thread thread1 = new Thread(printA);
10         Thread thread2 = new Thread(printB);
11         Thread thread3 = new Thread(print100);
12
13         // Start threads
14         thread1.start();
15         thread2.start();
16         thread3.start();
17     }
18 }
19
```



## Example: Using the Runnable Interface to Create and Launch Threads (2 of 3)

```
20 // The task for printing a specified character in specified times
21 class PrintChar implements Runnable {
22     private char charToPrint; // The character to print
23     private int times; // The times to repeat
24
25     /** Construct a task with specified character and number of
26      * times to print the character
27      */
28     public PrintChar(char c, int t) {
29         charToPrint = c;
30         times = t;
31     }
32
33     @Override /** Override the run() method to tell the system
34      * what the task to perform
35      */
36     public void run() {
37         for (int i = 0; i < times; i++) {
38             System.out.print(charToPrint);
39         }
40     }
41 }
42
```

## Example: Using the Runnable Interface to Create and Launch Threads (3 of 3)

```
43 // The task class for printing number from 1 to n for a given n
44 class PrintNum implements Runnable {
45     private int lastNum;
46
47     /** Construct a task for printing 1, 2, ... i */
48     public PrintNum(int n) {
49         lastNum = n;
50     }
51
52     @Override /** Tell the thread how to run */
53     public void run() {
54         for (int i = 1; i <= lastNum; i++) {
55             System.out.print(" " + i);
56         }
57     }
58 }
```

# The Thread Class

«interface»  
*java.lang.Runnable*



java.lang.Thread

+Thread()

Creates a default thread.

+Thread(task: Runnable)

Creates a thread for a specified task.

+start(): void

Starts the thread that causes the run() method to be invoked by the JVM.

+isAlive(): boolean

Tests whether the thread is currently running.

+setPriority(p: int): void

Sets priority p (ranging from 1 to 10) for this thread.

+join(): void

Waits for this thread to finish.

+sleep(millis: long): void

Puts the runnable object to sleep for a specified time in milliseconds.

+yield(): void

Causes this thread to temporarily pause and allow other threads to execute.

+interrupt(): void

Interrupts this thread.

# The Static `yield()` Method

You can use the `yield()` method to temporarily release time for other threads. For example, suppose you modify the code in Lines 53-57 in `TaskThreadDemo.java` as follows:

```
public void run() {  
    for (int i = 1; i <= lastNum; i++) {  
        System.out.print(" " + i);  
        Thread.yield();  
    }  
}
```

Every time a number is printed, the `print100` thread is yielded. So, the numbers are printed after the characters.

# The Static sleep(milliseconds) Method

The `sleep(long mills)` method puts the thread to sleep for the specified time in milliseconds. For example, suppose you modify the code in Lines 53-57 in `TaskThreadDemo.java` as follows:

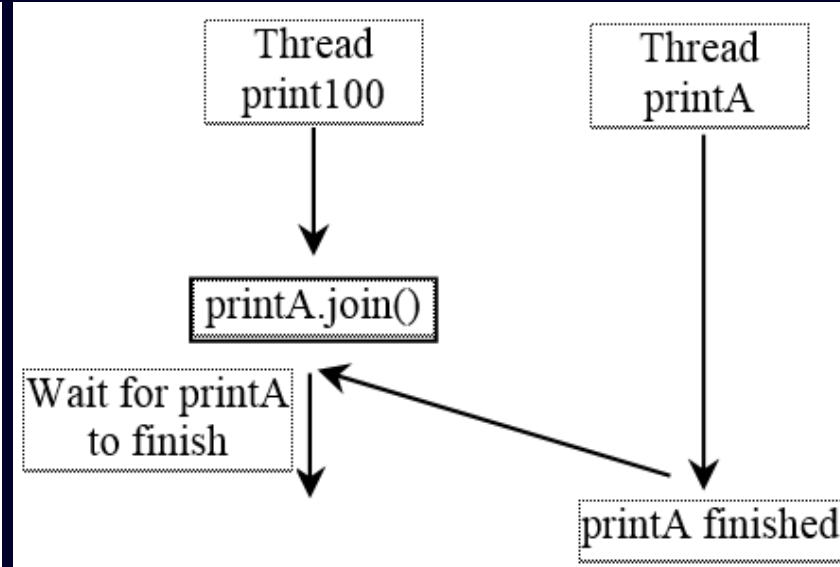
```
public void run() {  
    for (int i = 1; i <= lastNum; i++) {  
        System.out.print(" " + i);  
        try {  
            if (i >= 50) Thread.sleep(1);  
        }  
        catch (InterruptedException ex) {  
        }  
    }  
}
```

Every time a number ( $\geq 50$ ) is printed, the print100 thread is put to sleep for 1 millisecond.

# The () Method

You can use the **join()** method to force one thread to wait for another thread to finish. For example, suppose you modify the code in Lines 53-57 in `TaskThreadDemo.java` as follows:

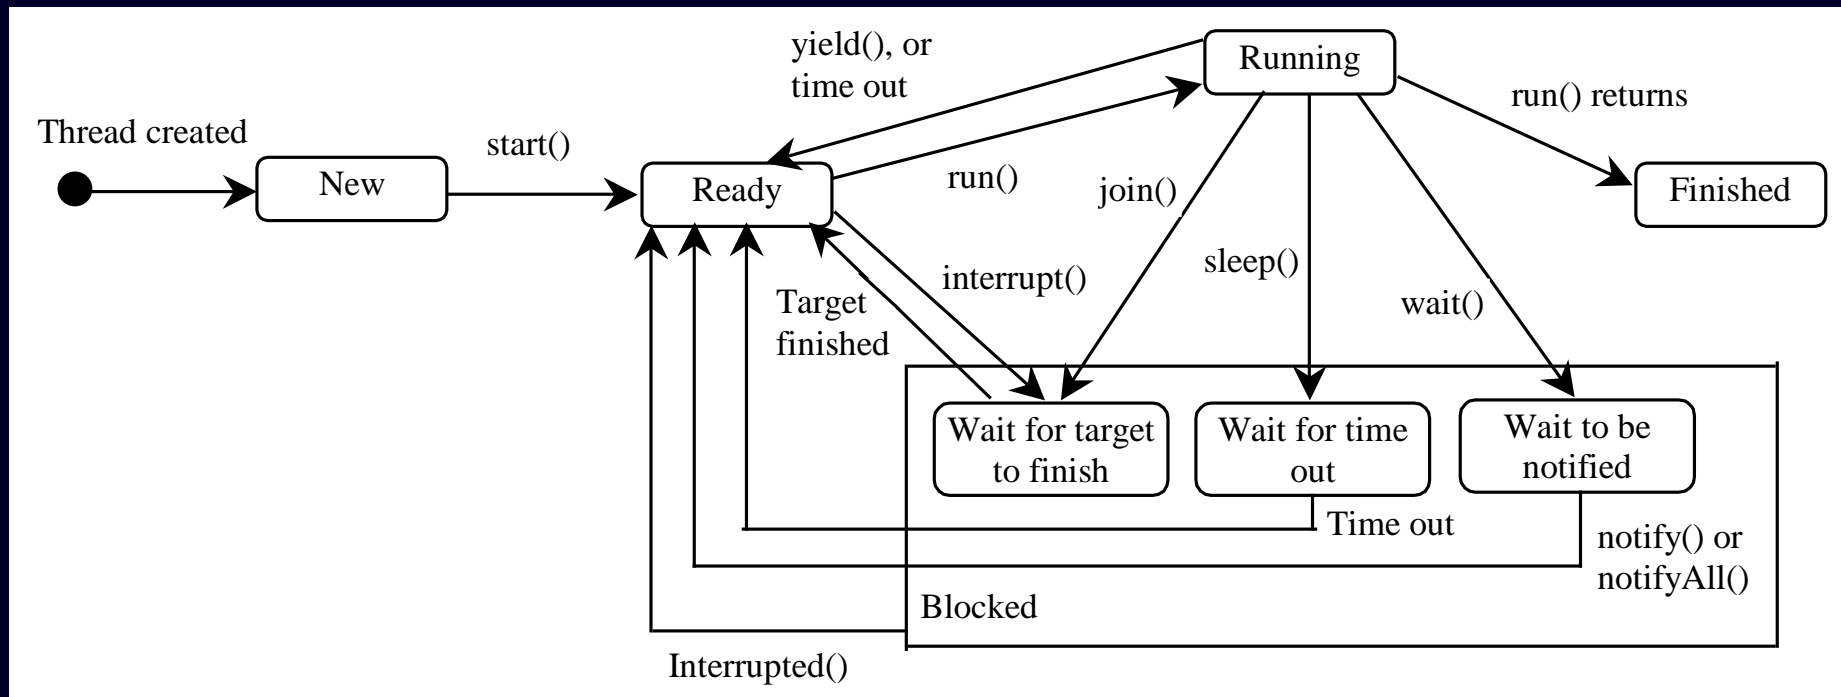
```
public void run() {  
    Thread threadA = new Thread( new PrintChar('c', 40));  
    threadA.start();  
  
    try {  
        for (int i = 1; i <= lastNum; i++) {  
            System.out.print(" " + i);  
  
            if (i == 50)  
                threadA.join();  
        }  
    } catch (InterruptedException ex) {  
    }  
}
```



The numbers after 50 are printed after thread `printA` is finished.

# Thread States

A thread can be in one of five states: **New**, **Ready**, **Running**, **Blocked**, or **Finished**.



# isAlive(), interrupt(), and isInterrupted()

The **isAlive()** method is used to find out the state of a thread. It returns **true** if a thread is in the **Ready, Blocked, or Running state**; it returns false if a thread is **new and has not started or if it is finished**.

The **interrupt()** method interrupts a thread in the following way: If a thread is currently in the Ready or Running state, its interrupted flag is set; if a thread is currently blocked, it is awakened and enters the Ready state, and an `java.io.InterruptedException` is thrown.

The **isInterrupted()** method tests whether the thread is interrupted.



# The deprecated stop(), suspend(), and resume() Methods

NOTE: The **Thread** class also contains the **stop()**, **suspend()**, and **resume()** methods. As of Java 2, these methods are *deprecated* (or *outdated*) because they are known to be inherently unsafe. You should assign **null** to a **Thread** variable to indicate that it is stopped rather than use the **stop()** method.

# Thread Priority

- Each thread is assigned a default priority of `Thread.NORM_PRIORITY`. You can reset the priority using `setPriority(int priority)`.
- Some constants for priorities include  
`Thread.MIN_PRIORITY`  
`Thread.MAX_PRIORITY`  
`Thread.NORM_PRIORITY`

# Example: Flashing Text

```
import javax.swing.*;

public class FlashingText extends JApplet implements Runnable {
    private JLabel jlblText = new JLabel("Welcome", JLabel.CENTER);

    public FlashingText() {
        add(jlblText);
        new Thread(this).start();
    }

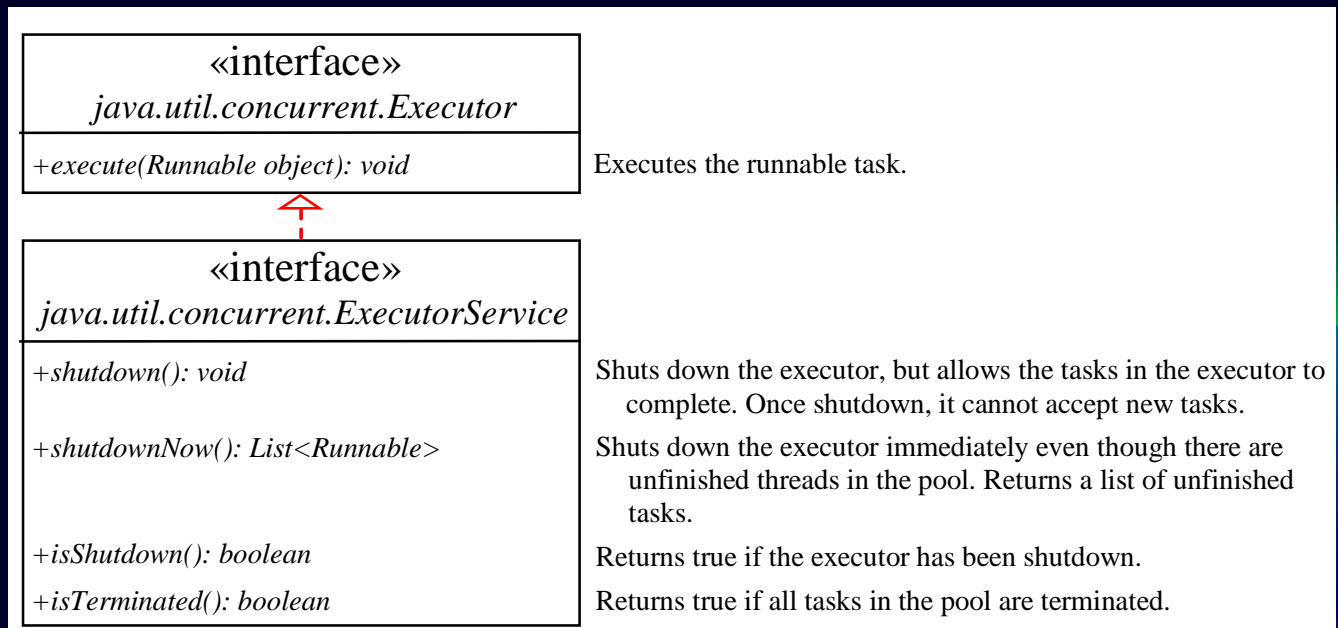
    @Override /** Set the text on/off every 200 milliseconds */
    public void run() {
        try {
            while (true) {
                if (jlblText.getText() == null)
                    jlblText.setText("Welcome");
                else
                    jlblText.setText(null);

                Thread.sleep(200);
            }
        } catch (InterruptedException ex) {
        }
    }

    /** Main method */
    public static void main(String[] args) {
        SwingUtilities.invokeLater(new Runnable() {
            public void run() {
                JFrame frame = new JFrame("FlashingText");
                frame.add(new FlashingText());
                frame.setLocationRelativeTo(null); // Center the frame
                frame.setDefaultCloseOperation(JFrame.EXIT_ON_CLOSE);
                frame.setSize(200, 200);
                frame.setVisible(true);
            }
        });
    }
}
```

# Thread Pools

Starting a new thread for each task could limit throughput and cause poor performance. A thread pool is ideal to manage the number of tasks executing concurrently. Starting with JDK 1.5 you can use the Executor interface for executing tasks in a thread pool and the ExecutorService interface for managing and controlling tasks. ExecutorService is a subinterface of Executor.



# Creating Executors

To create an **Executor** object, use the static methods in the **Executors** class.

java.util.concurrent.Executors
+newFixedThreadPool(numberOfThreads: int): ExecutorService
+newCachedThreadPool(): ExecutorService

Creates a thread pool with a fixed number of threads executing concurrently. A thread may be reused to execute another task after its current task is finished.

Creates a thread pool that creates new threads as needed, but will reuse previously constructed threads when they are available.

# Creating Executors

```
1  import java.util.concurrent.*;
2
3  public class ExecutorDemo {
4      public static void main(String[] args) {
5          // Create a fixed thread pool with maximum three threads
6          ExecutorService executor = Executors.newFixedThreadPool(3);
7
8          // Submit runnable tasks to the executor
9          executor.execute(new PrintChar('a', 100));
10         executor.execute(new PrintChar('b', 100));
11         executor.execute(new PrintNum(100));
12
13         // Shut down the executor
14         executor.shutdown();
15     }
16 }
```

# Thread Synchronization

A shared resource may be corrupted if it is accessed simultaneously by multiple threads. For example, two unsynchronized threads accessing the same bank account may cause conflict.

Step	balance	thread[i]	thread[j]
1	0	<code>newBalance = bank.getBalance() + 1;</code>	
2	0		<code>newBalance = bank.getBalance() + 1;</code>
3	1	<code>bank.setBalance(newBalance);</code>	
4	1		<code>bank.setBalance(newBalance);</code>

# Race Condition

What caused the error in the example? Here is a possible scenario:

Step	balance	Task 1	Task 2
1	0	newBalance = balance + 1;	
2	0		newBalance = balance + 1;
3	1	balance = newBalance;	
4	1		balance = newBalance;

- The effect of this scenario is that **Task 1** did nothing, because in **Step 4 Task 2** overrides **Task 1's** result. Obviously, the problem is that **Task 1** and **Task 2** are accessing a common resource in a way that causes conflict. This is a common problem known as a *race condition* in multithreaded programs.
- A class is said to be *thread-safe* if an object of the class does not cause a race condition in the presence of multiple threads.



# The synchronized keyword

- To avoid race conditions, more than one thread must be prevented from simultaneously entering certain part of the program, known as **critical region (critical section)**.
- You can use the **synchronized** keyword to synchronize the method so that only one thread can access the method at a time.

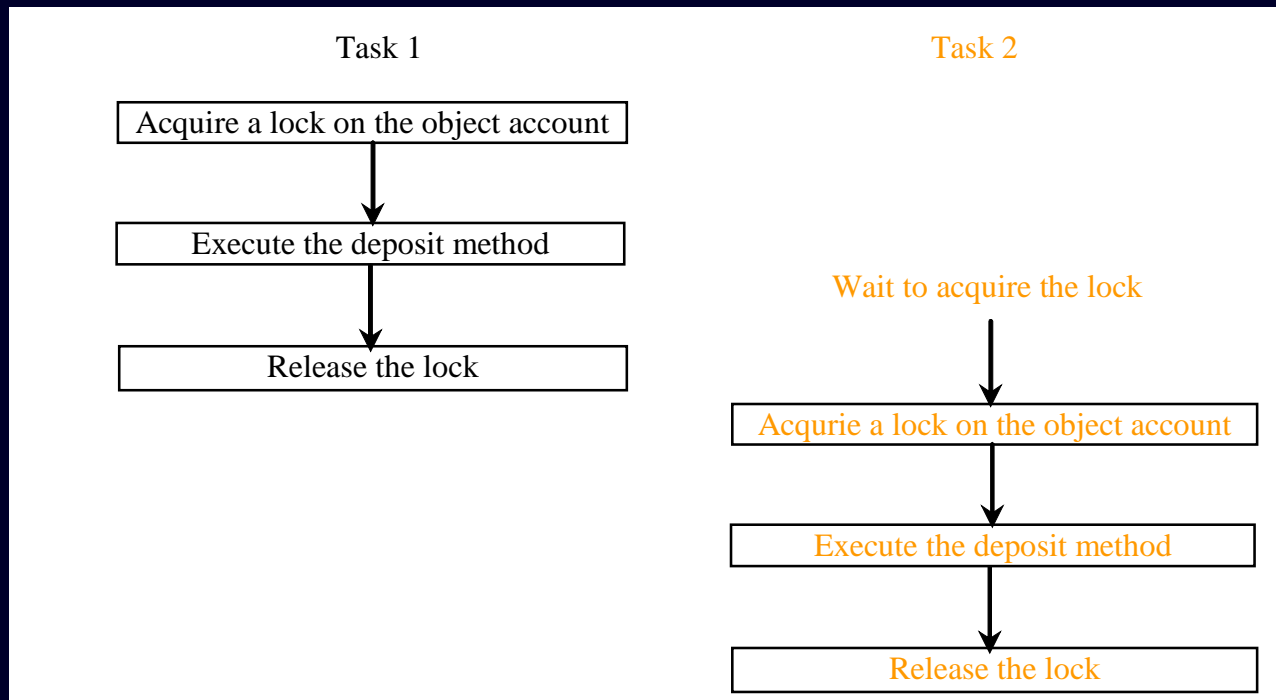
```
public synchronized void withdraw(double amount)
```

# Synchronizing Instance Methods and Static Methods

- A synchronized method acquires a **lock** before it executes.
- In the case of an instance method, the lock is on the object for which the method was invoked. In the case of a static method, the lock is on the class.
- If one thread invokes a synchronized instance method (respectively, static method) on an object, the lock of that object (respectively, class) is acquired first, then the method is executed, and finally the lock is released.
- Another thread invoking the same method of that object (respectively, class) is blocked until the lock is released.

# Synchronizing Instance Methods and Static Methods

With the deposit method synchronized, the preceding scenario cannot happen. If **Task 2** starts to enter the method, and **Task 1** is already in the method, **Task 2** is blocked until **Task 1** finishes the method.



# Synchronizing Statements

Invoking a synchronized instance method of an object acquires a lock on the object, and invoking a synchronized static method of a class acquires a lock on the class.

A synchronized statement can be used to acquire a lock on any object, not just *this* object, when executing a block of the code in a method. This block is referred to as a *synchronized block*. The general form of a synchronized statement is as follows:

```
synchronized (expr) {  
    statements;  
}
```

The expression `expr` must evaluate to an object reference. If the object is already locked by another thread, the thread is blocked until the lock is released. When a lock is obtained on the object, the statements in the synchronized block are executed, and then the lock is released.

# Synchronizing Statements vs. Methods

Any synchronized instance method can be converted into a synchronized statement. Suppose that the following is a synchronized instance method:

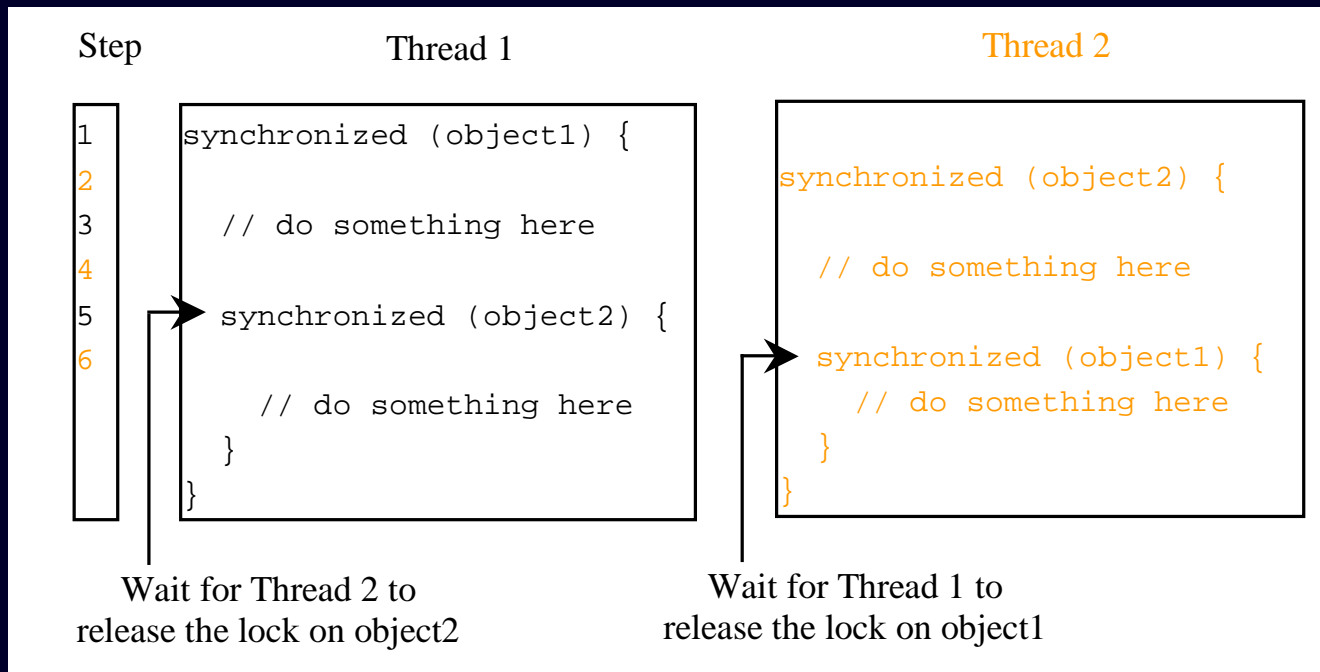
```
public synchronized void xMethod() {  
    // method body  
}
```

This method is equivalent to

```
public void xMethod() {  
    synchronized (this) {  
        // method body  
    }  
}
```

# Deadlock

Sometimes two or more threads need to acquire the locks on several shared objects. This could cause *deadlock*, in which each thread has the lock on one of the objects and is waiting for the lock on the other object. Consider the scenario with two threads and two objects, as shown below. **Thread 1** acquired a lock on **object1** and **Thread 2** acquired a lock on **object2**. Now Thread 1 is waiting for the lock on **object2** and **Thread 2** for the lock on **object1**. The two threads wait for each other to release the in order to get the lock, and neither can continue to run.



# Preventing Deadlock

- Deadlock can be easily avoided by using a simple technique known as resource ordering.
- With this technique, you assign an order on all the objects whose locks must be acquired and ensure that each thread acquires the locks in that order.
- For the example in the previous figure, suppose the objects are ordered as **object1** and **object2**.
- Using the resource ordering technique, **Thread 2** must acquire a lock on **object1** first, then on **object2**. Once **Thread 1** acquired a lock on **object1**, **Thread 2** has to wait for a lock on **object1**.
- So **Thread 1** will be able to acquire a lock on **object2** and no deadlock would occur.

# Synchronized Collections

The classes in the Java Collections Framework are not thread-safe, i.e., the contents may be corrupted if they are accessed and updated concurrently by multiple threads. You can protect the data in a collection by locking the collection or using synchronized collections.

The **Collections** class provides six static methods for wrapping a collection into a synchronized version. The collections created using these methods are called *synchronization wrappers*.

## java.util.Collections

+synchronizedCollection(c: Collection): Collection

Returns a synchronized collection.

+synchronizedList(list: List): List

Returns a synchronized list from the specified list.

+synchronizedMap(m: Map): Map

Returns a synchronized map from the specified map.

+synchronizedSet(s: Set): Set

Returns a synchronized set from the specified set.

+synchronizedSortedMap(s: SortedMap): SortedMap

Returns a synchronized sorted map from the specified sorted map.

+synchronizedSortedSet(s: SortedSet): SortedSet

Returns a synchronized sorted set.



# Vector, Stack, and Hashtable

Invoking `synchronizedCollection(Collection c)` returns a new `Collection` object, in which all the methods that access and update the original collection `c` are synchronized. These methods are implemented using the `synchronized` keyword. For example, the `add` method is implemented like this:

```
public boolean add(E o) {  
    synchronized (this) {  
        return c.add(o);  
    }  
}
```

The synchronized collections can be safely accessed and modified by multiple threads concurrently.

The methods in `java.util.Vector`, `java.util.Stack`, and `Hashtable` are already synchronized. These are old classes introduced in JDK 1.0. Starting with JDK 1.5, you should use `java.util.ArrayList` to replace `Vector`, `java.util.LinkedList` to replace `Stack`, and `java.util.Map` to replace `Hashtable`. If synchronization is needed, use a synchronization wrapper.

# Programming Challenge

Write a program that launches 1,000 threads. Each thread adds **1** to a variable **sum** that initially is **0**.

Define an **int** variable to hold **sum**.

Run the program with and without synchronization to see its effect.

